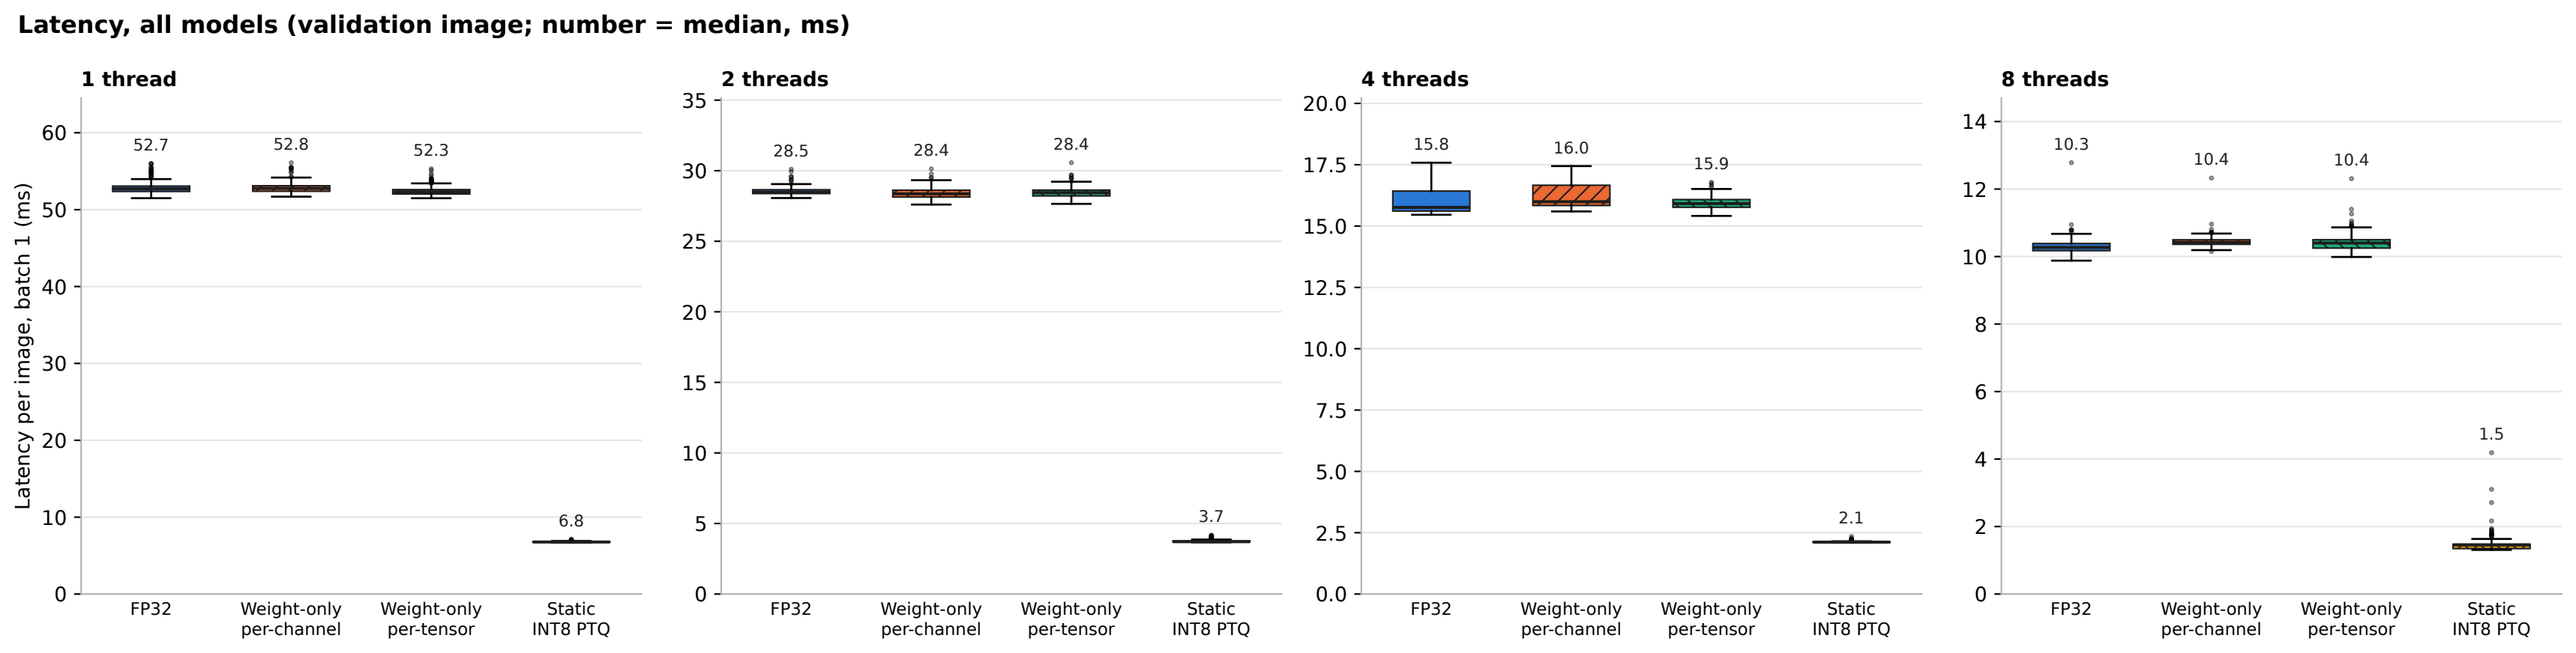
Latency, all models (validation image; number = median, ms)
1 thread
2 threads
4 threads
8 threads
Latency per image, batch 1 (ms)
52.7
52.8
52.3
6.8
28.5
28.4
28.4
3.7
15.8
16.0
15.9
2.1
10.3
10.4
10.4
1.5
FP32
Weight-only per-channel
Weight-only per-tensor
Static INT8 PTQ