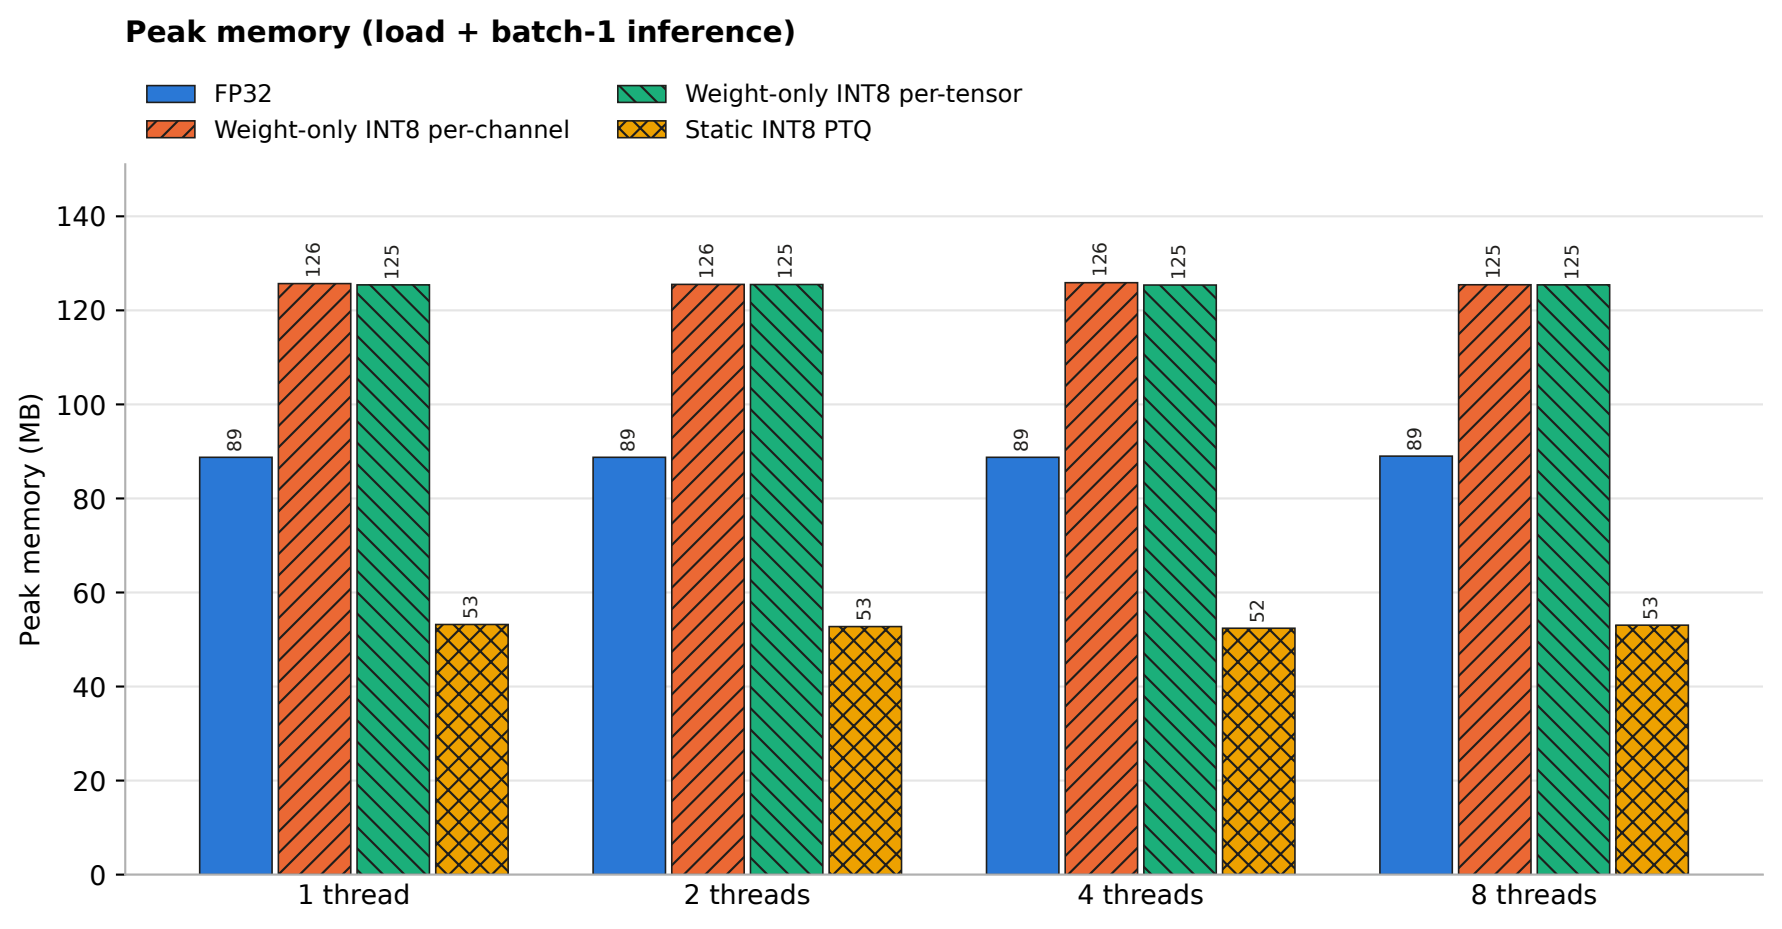
Peak memory (load + batch-1 inference)
FP32
Weight-only INT8 per-channel
Weight-only INT8 per-tensor
Static INT8 PTQ
Peak memory (MB)
0
20
40
60
80
100
120
140
89
126
125
53
89
126
125
53
89
126
125
52
89
125
125
53
1 thread
2 threads
4 threads
8 threads