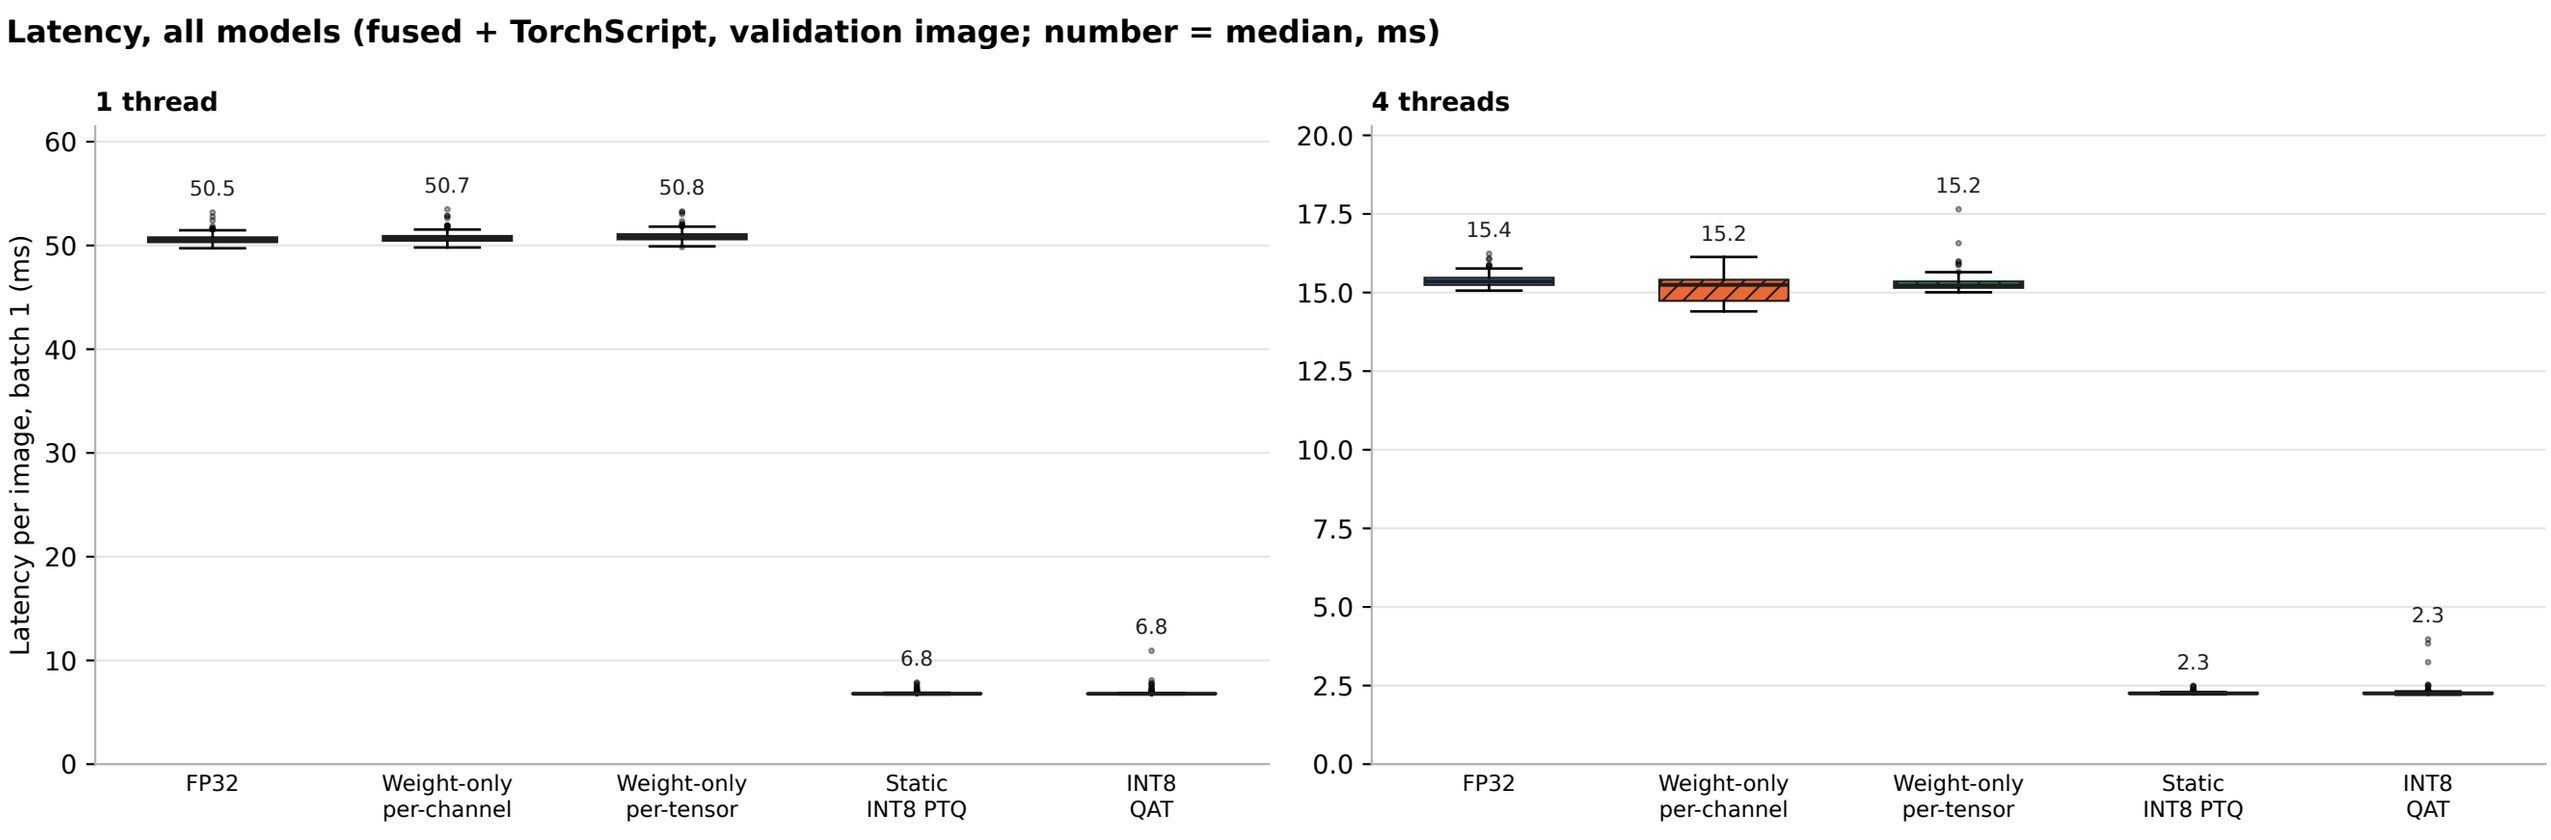
Latency, all models (fused + TorchScript, validation image; number = median, ms)
1 thread
4 threads
Latency per image, batch 1 (ms)
0
10
20
30
40
50
60
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
20.0
50.5
50.7
50.8
6.8
6.8
15.4
15.2
15.2
2.3
2.3
FP32
Weight-only per-channel
Weight-only per-tensor
Static INT8 PTQ
INT8 QAT