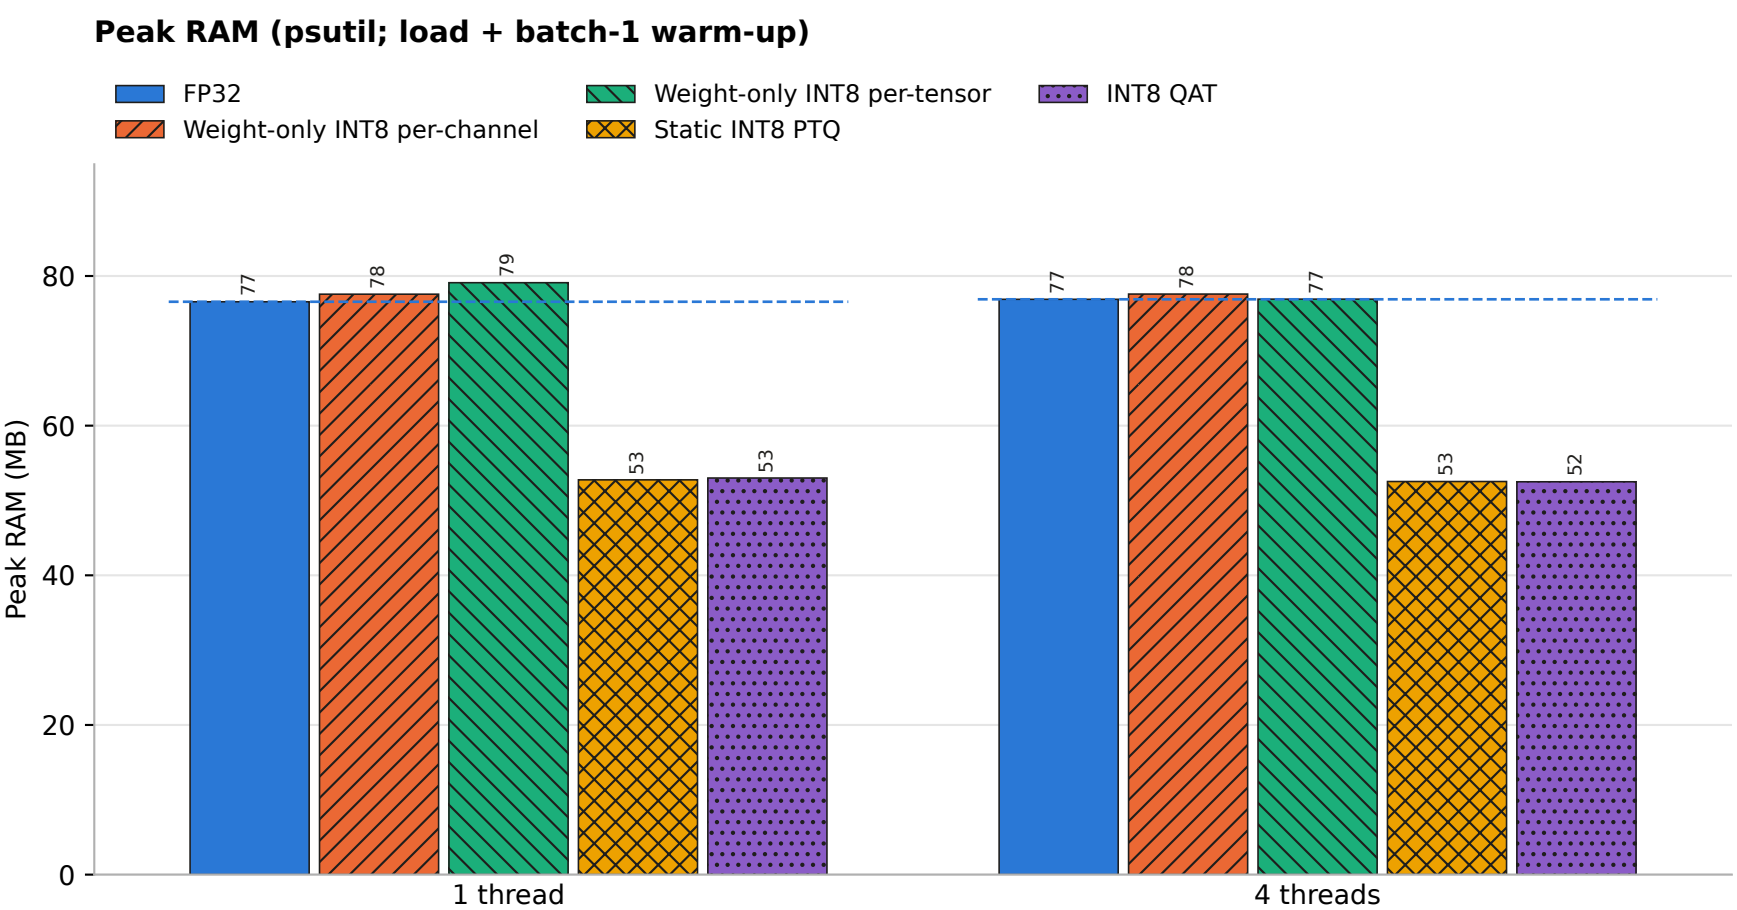
Peak RAM (psutil; load + batch-1 warm-up)
FP32
Weight-only INT8 per-channel
Weight-only INT8 per-tensor
Static INT8 PTQ
INT8 QAT
Peak RAM (MB)
0
20
40
60
80
77
78
79
53
53
77
78
77
53
52
1 thread
4 threads