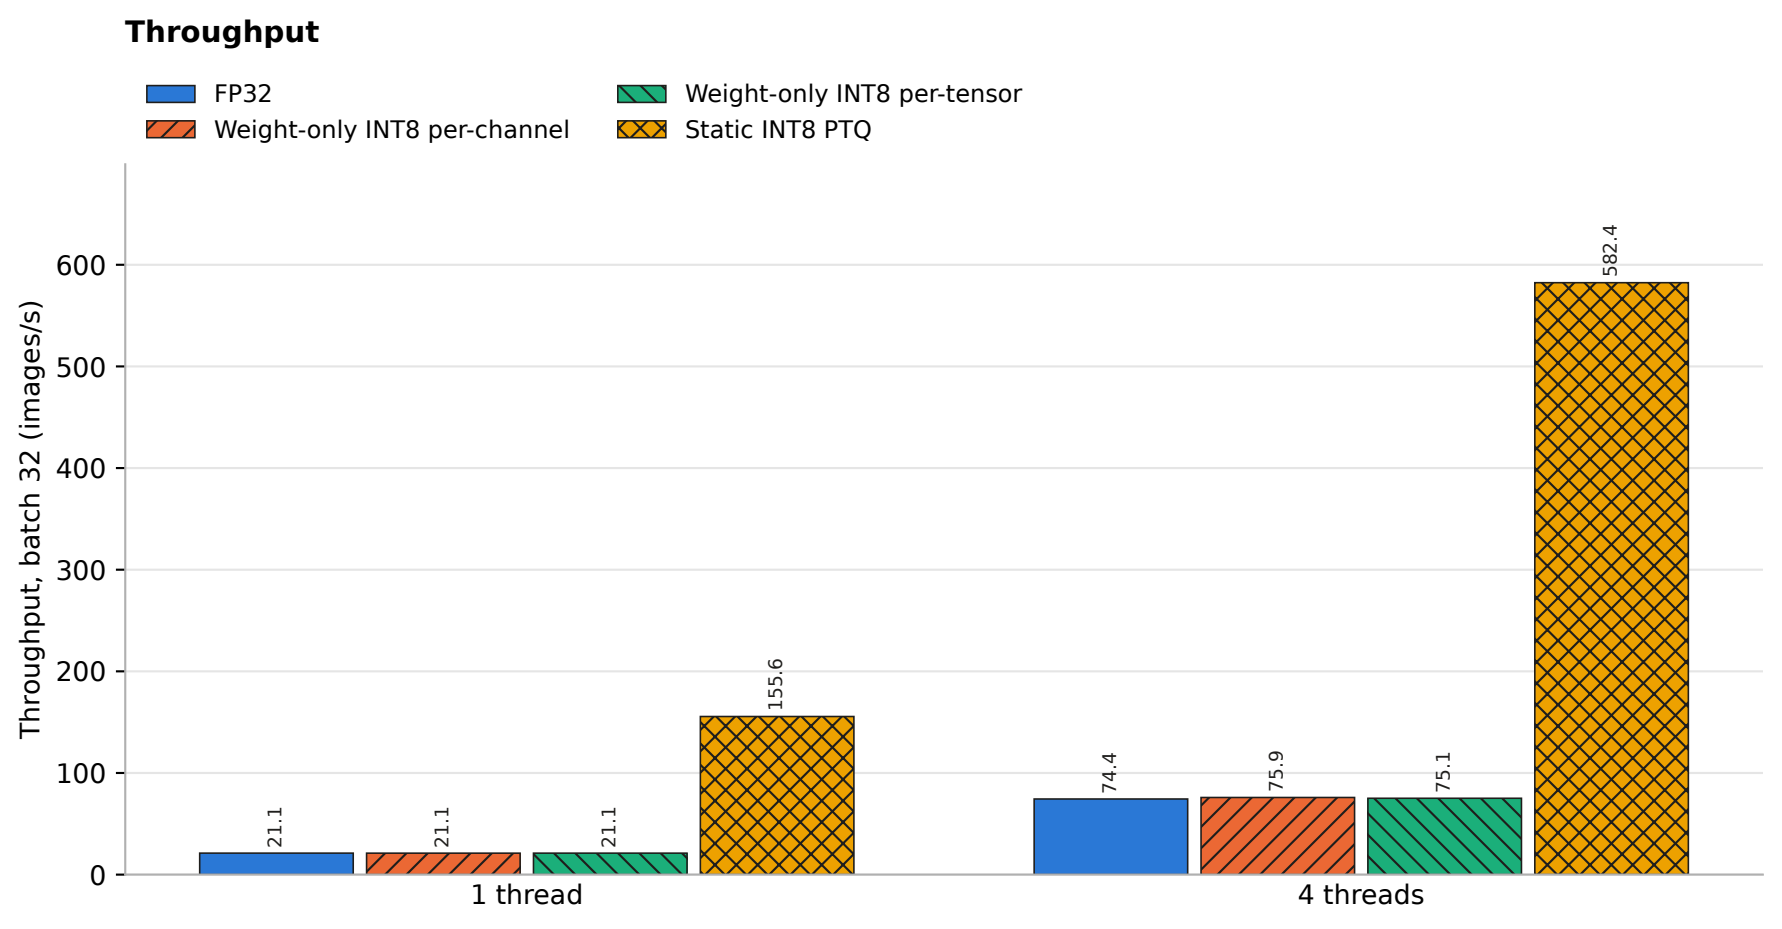
Throughput
FP32
Weight-only INT8 per-channel
Weight-only INT8 per-tensor
Static INT8 PTQ
Throughput, batch 32 (images/s)
0
100
200
300
400
500
600
21.1
21.1
21.1
155.6
74.4
75.9
75.1
582.4
1 thread
4 threads