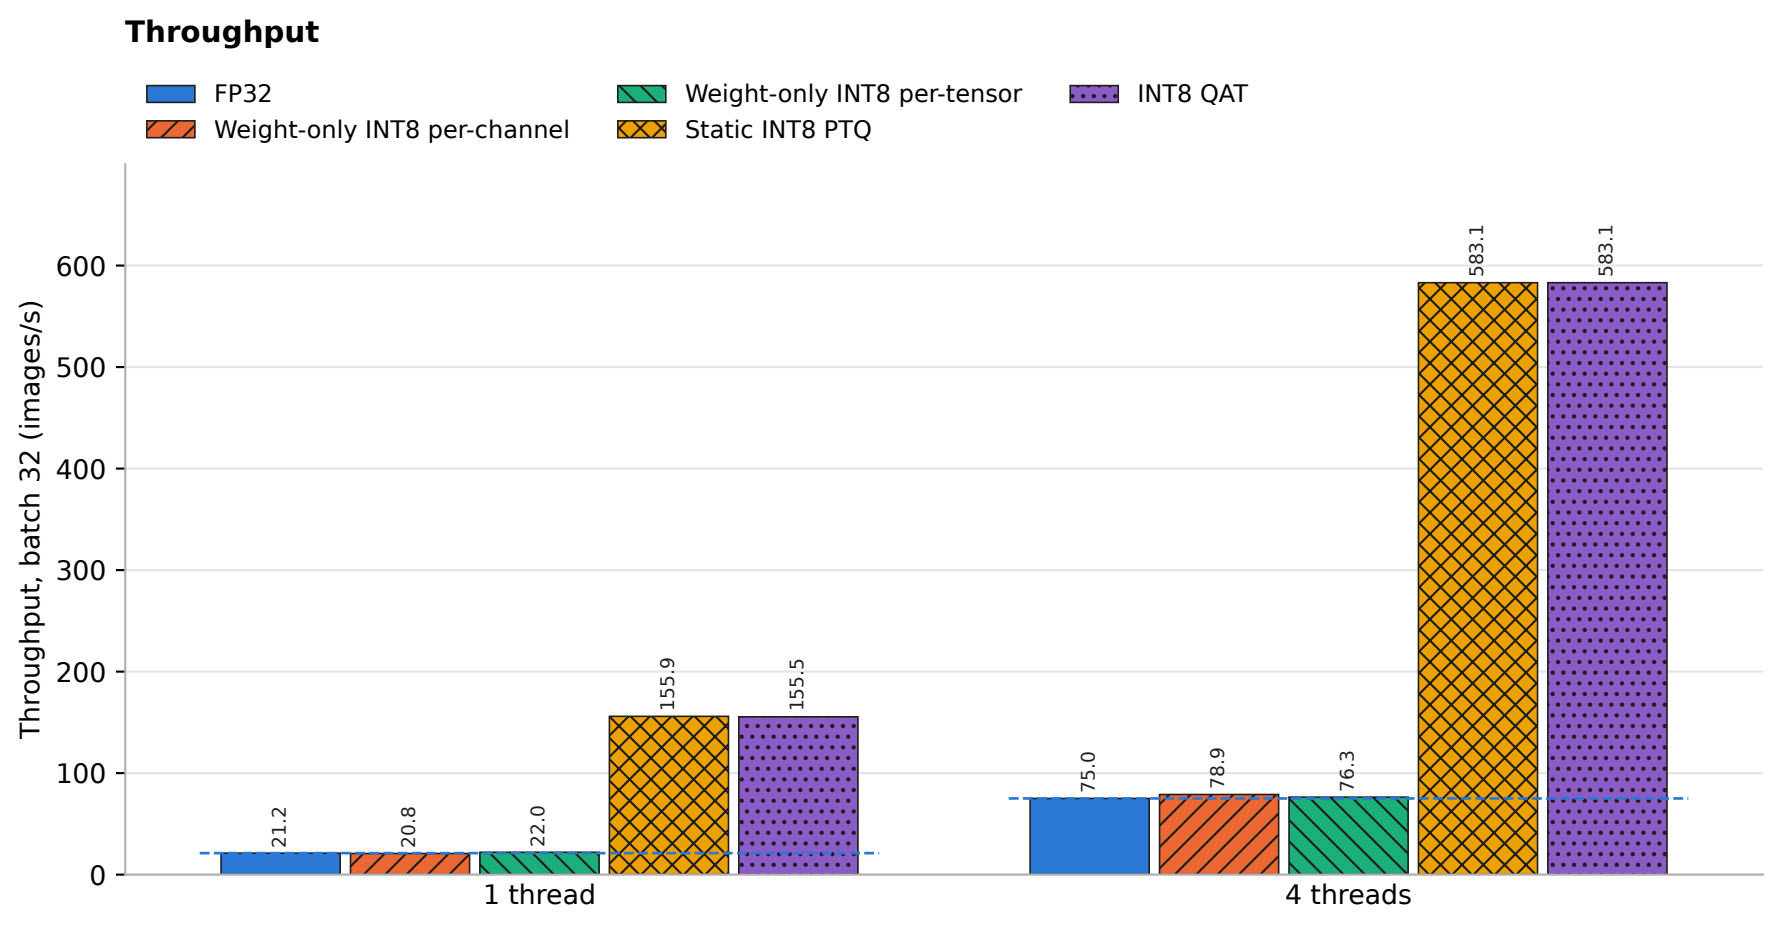
Throughput
FP32
Weight-only INT8 per-channel
Weight-only INT8 per-tensor
Static INT8 PTQ
INT8 QAT
Throughput, batch 32 (images/s)
0
100
200
300
400
500
600
21.2
20.8
22.0
155.9
155.5
75.0
78.9
76.3
583.1
583.1
1 thread
4 threads