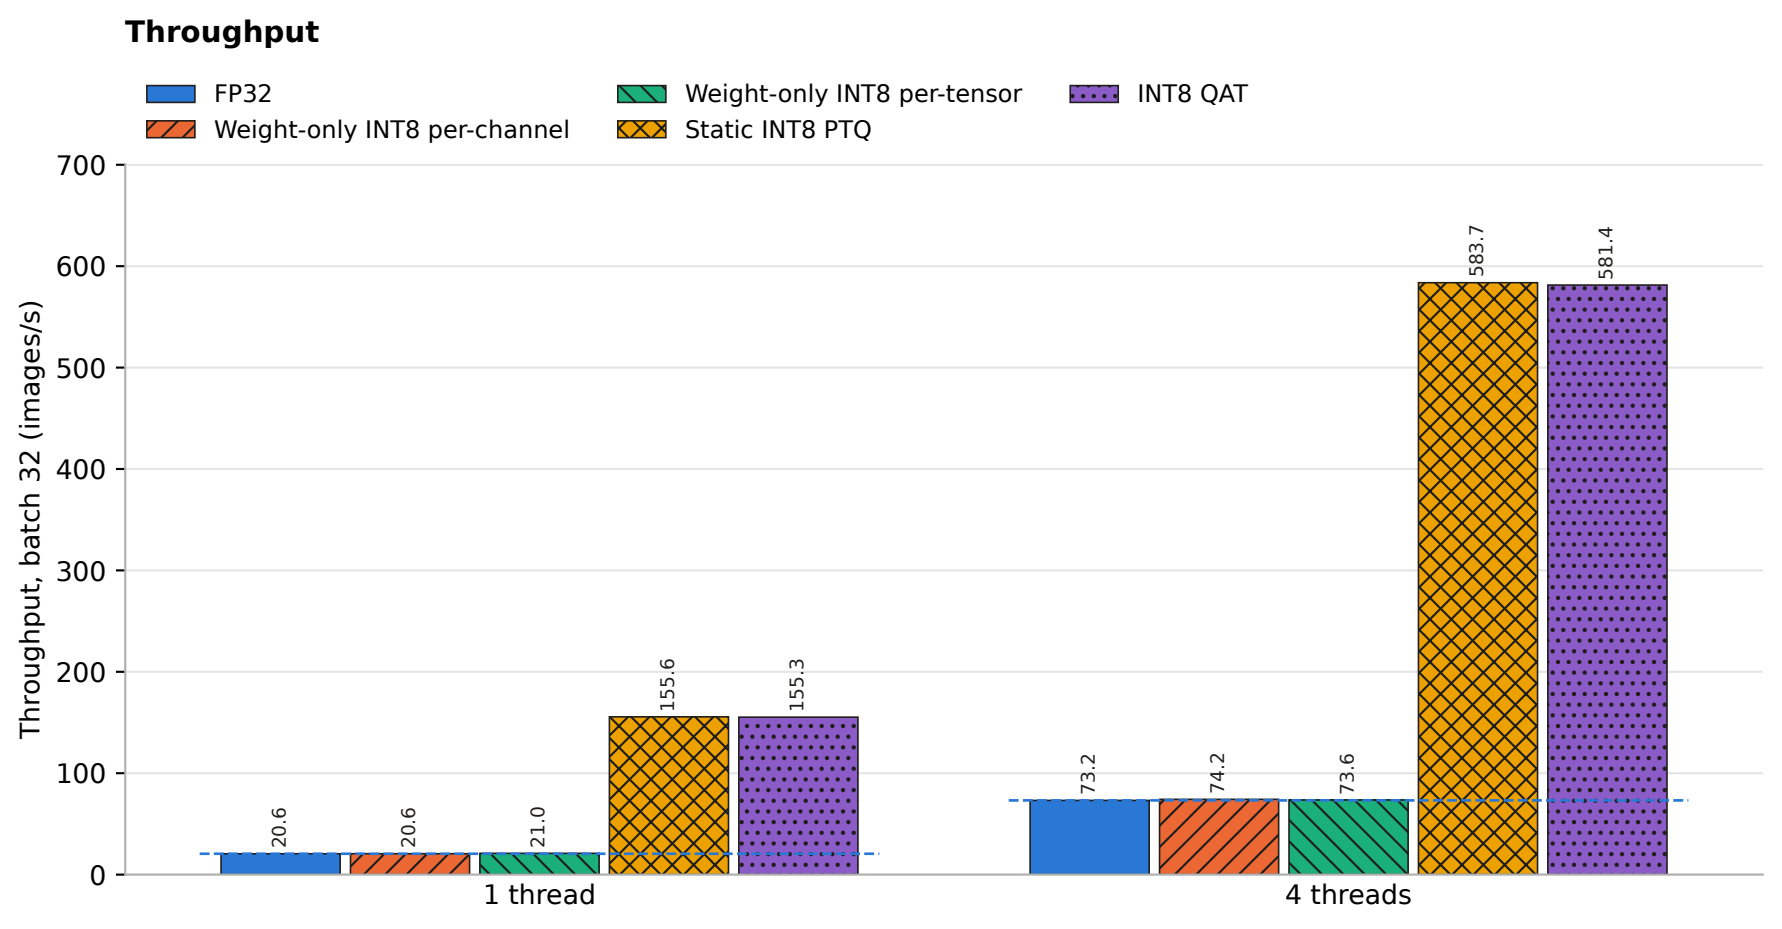
Throughput
FP32
Weight-only INT8 per-channel
Weight-only INT8 per-tensor
Static INT8 PTQ
INT8 QAT
Throughput, batch 32 (images/s)
700
600
500
400
300
200
100
0
20.6
20.6
21.0
155.6
155.3
73.2
74.2
73.6
583.7
581.4
1 thread
4 threads