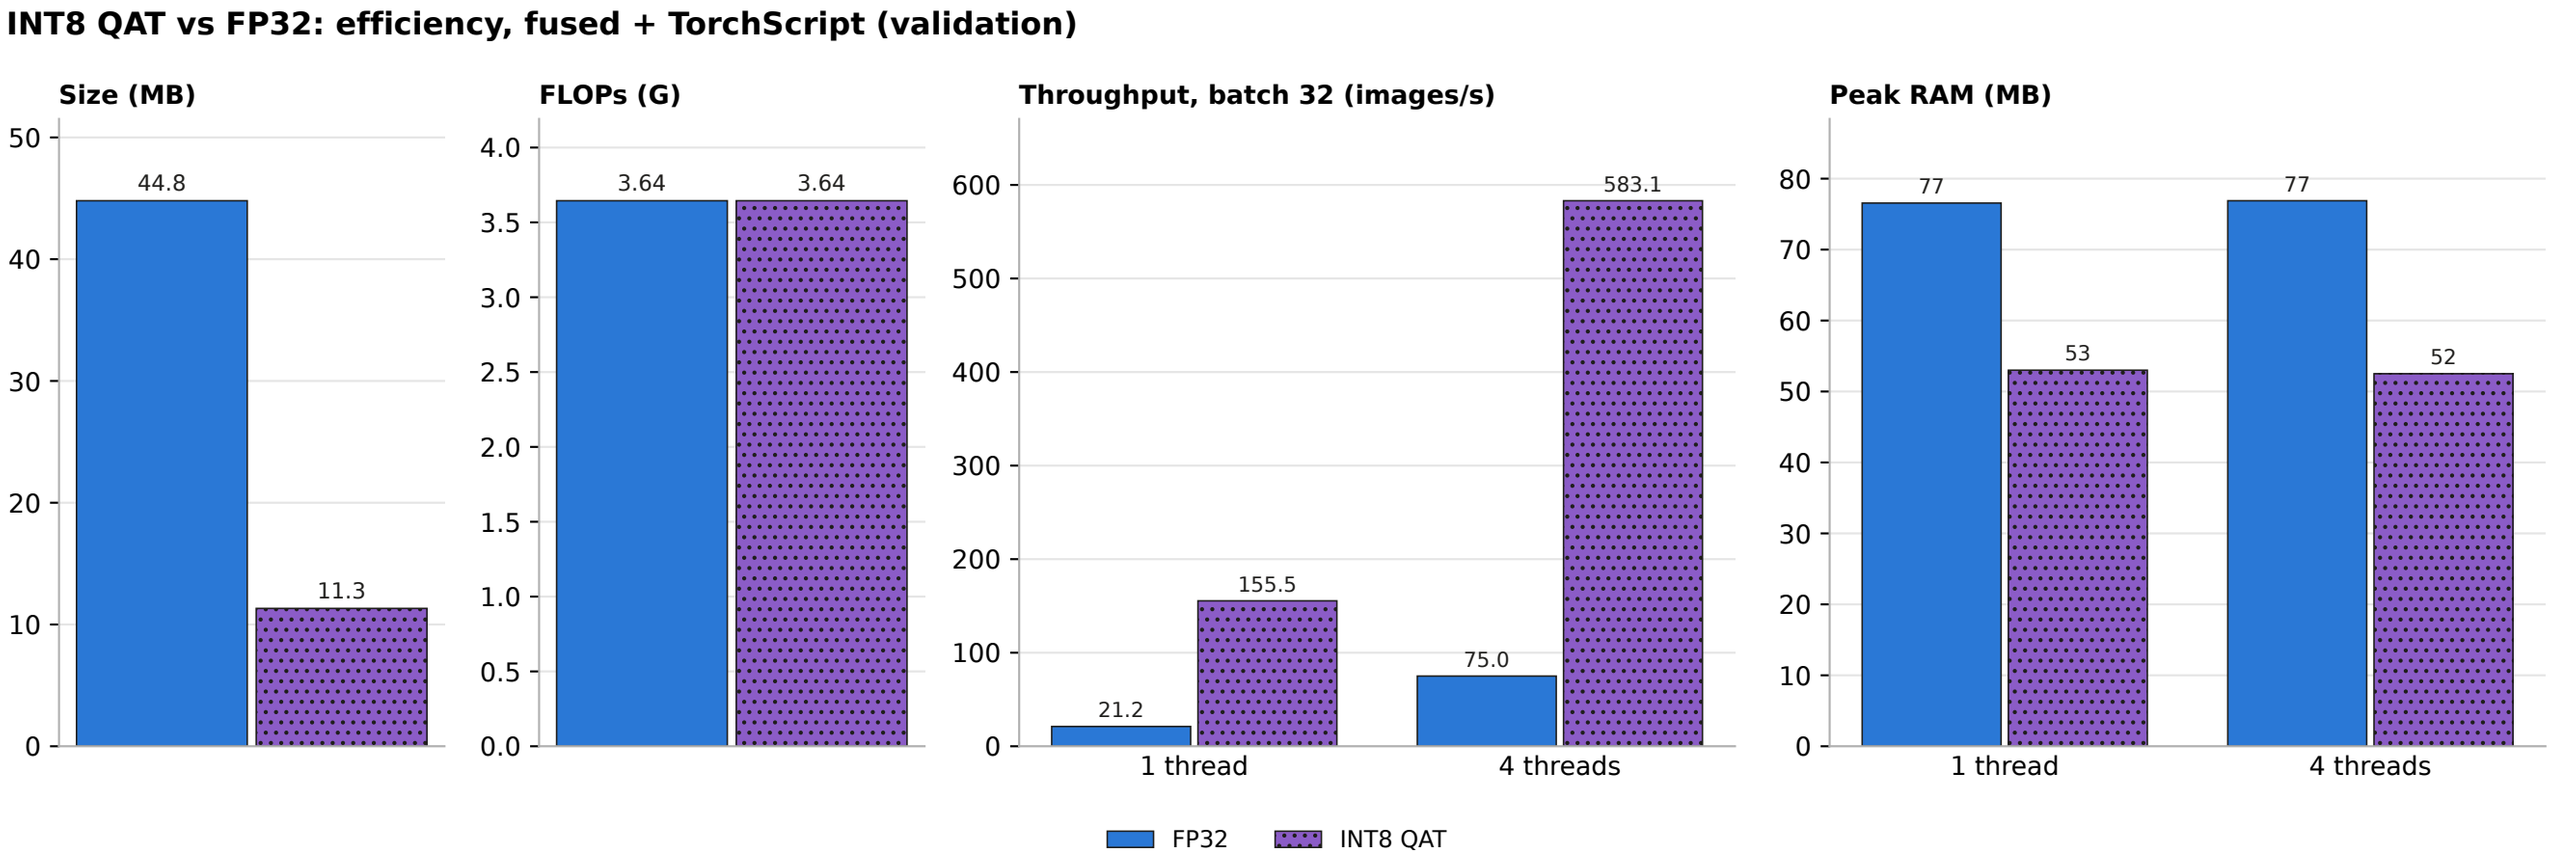
INT8 QAT vs FP32: efficiency, fused + TorchScript (validation)
Size (MB)
44.8
11.3
FLOPs (G)
3.64
3.64
Throughput, batch 32 (images/s)
21.2
155.5
75.0
583.1
1 thread
4 threads
Peak RAM (MB)
77
53
77
52
1 thread
4 threads
FP32
INT8 QAT