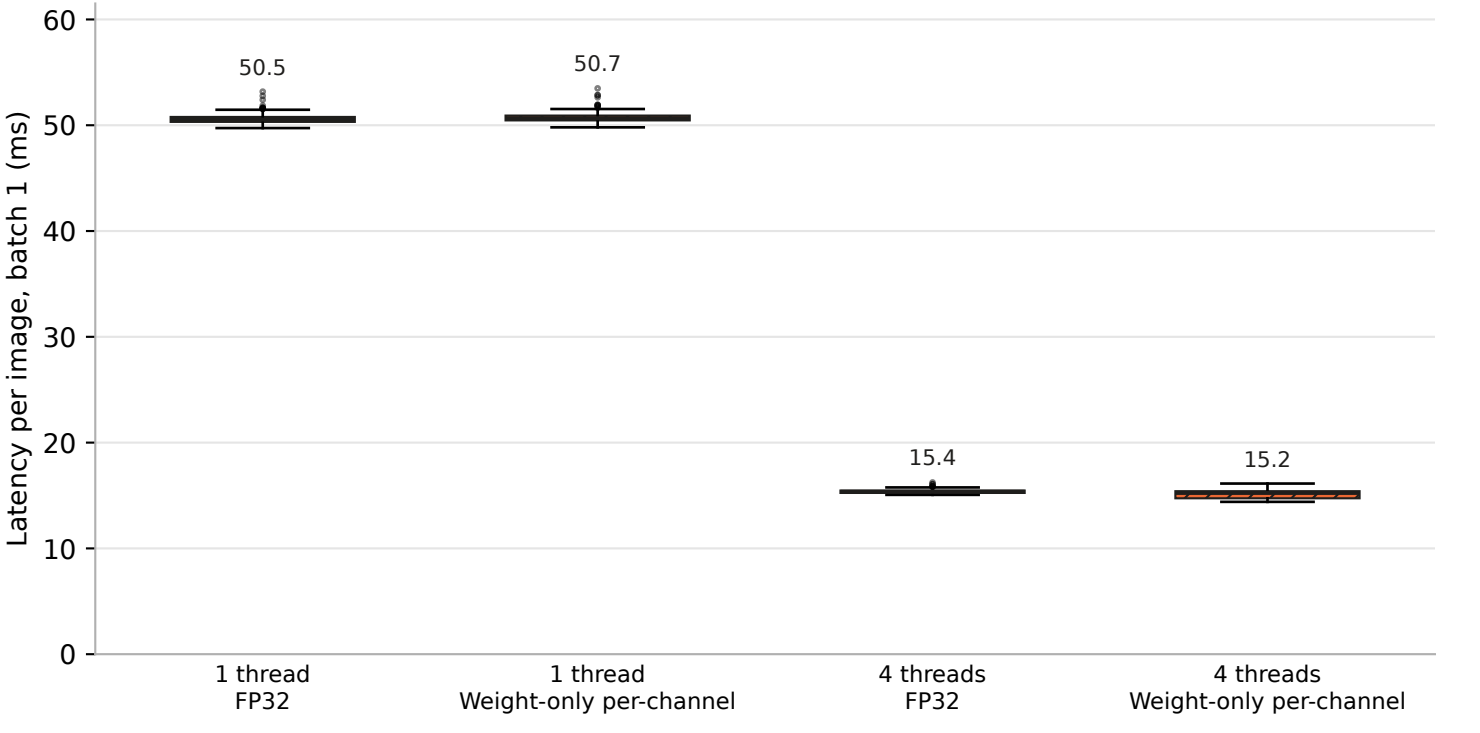
Weight-only INT8 per-channel vs FP32: latency (number = median, ms) (validation)
Latency per image, batch 1 (ms)
0
10
20
30
40
50
60
50.5
50.7
15.4
15.2
1 thread
FP32
1 thread
Weight-only per-channel
4 threads
FP32
4 threads
Weight-only per-channel
Each box: 600 timed runs (3 repeats x 200). Format: fused + TorchScript.