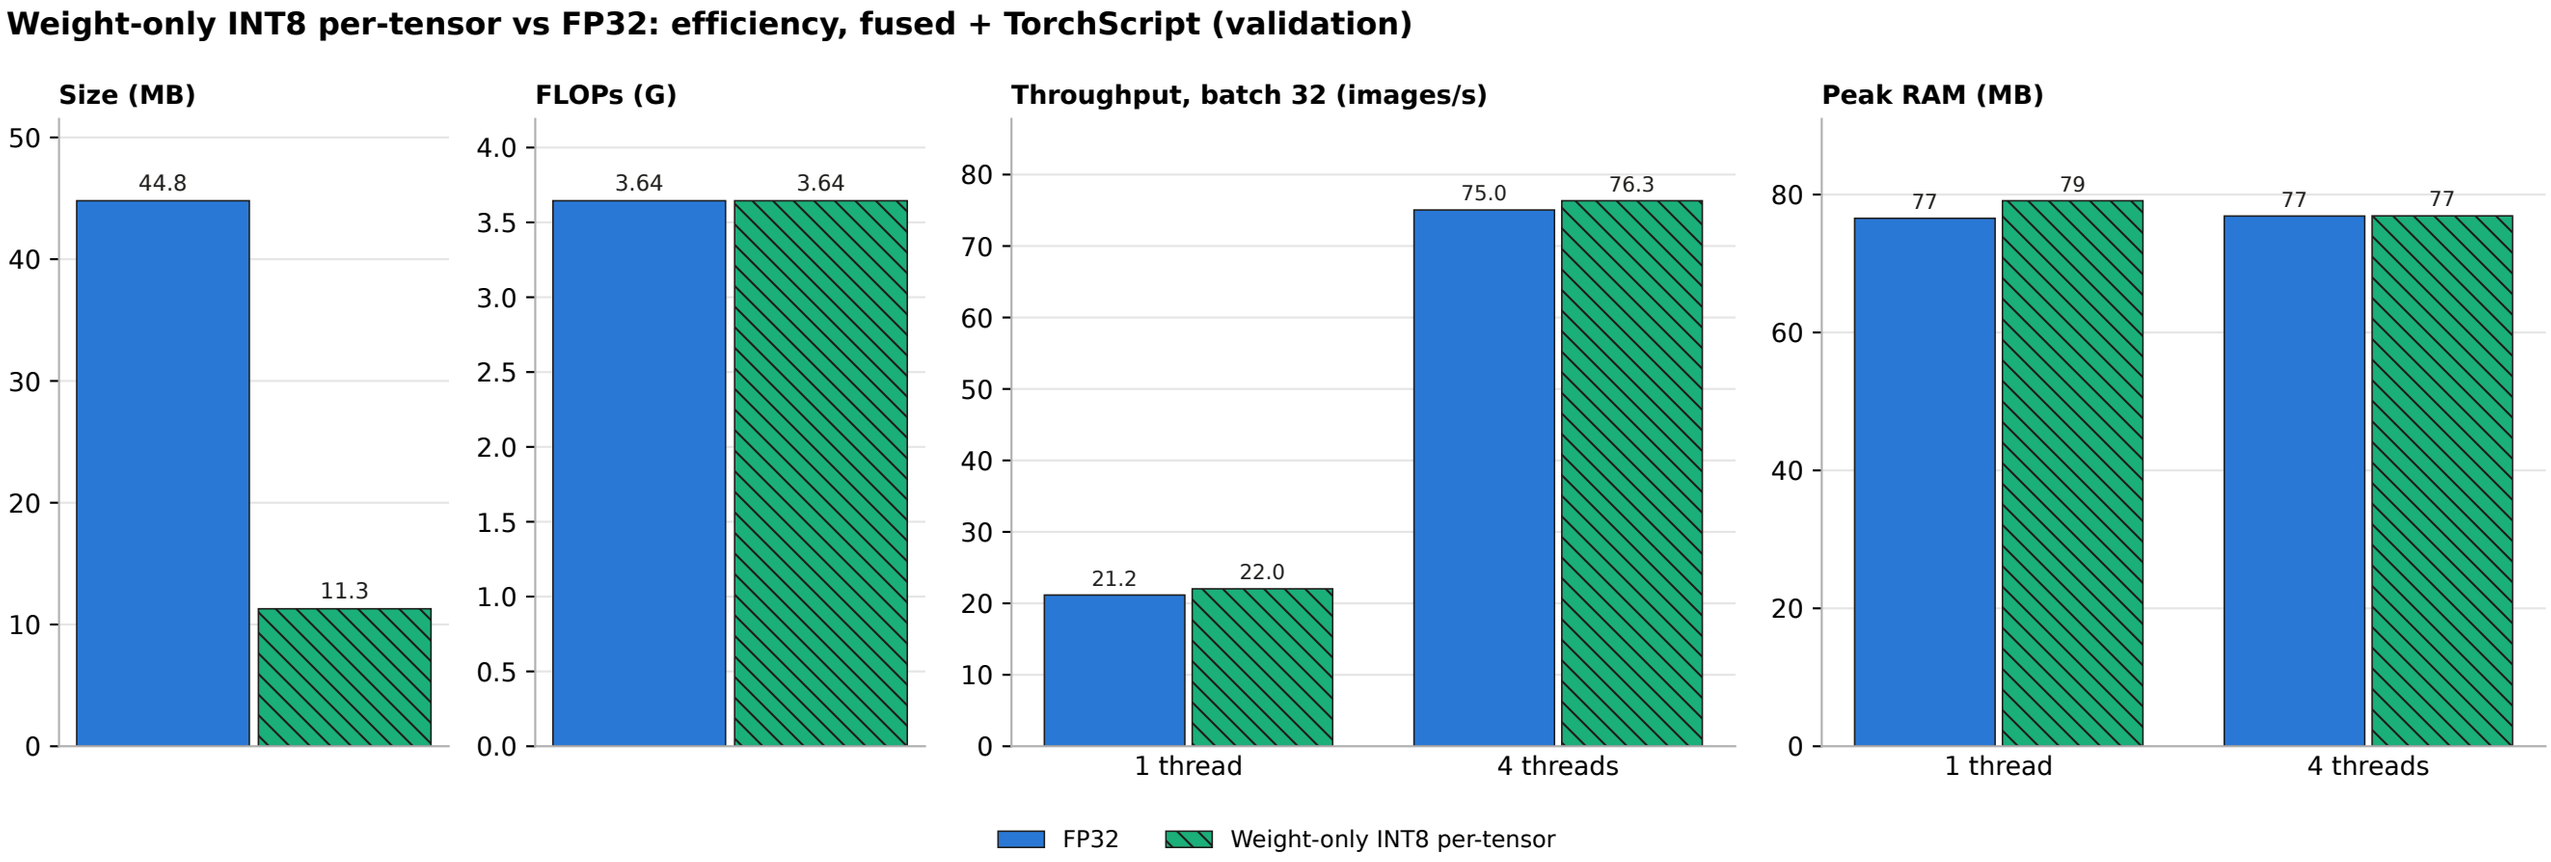
Weight-only INT8 per-tensor vs FP32: efficiency, fused + TorchScript (validation)
Size (MB)
44.8
11.3
FLOPs (G)
3.64
3.64
Throughput, batch 32 (images/s)
21.2
22.0
75.0
76.3
1 thread
4 threads
Peak RAM (MB)
77
79
77
77
1 thread
4 threads
FP32
Weight-only INT8 per-tensor