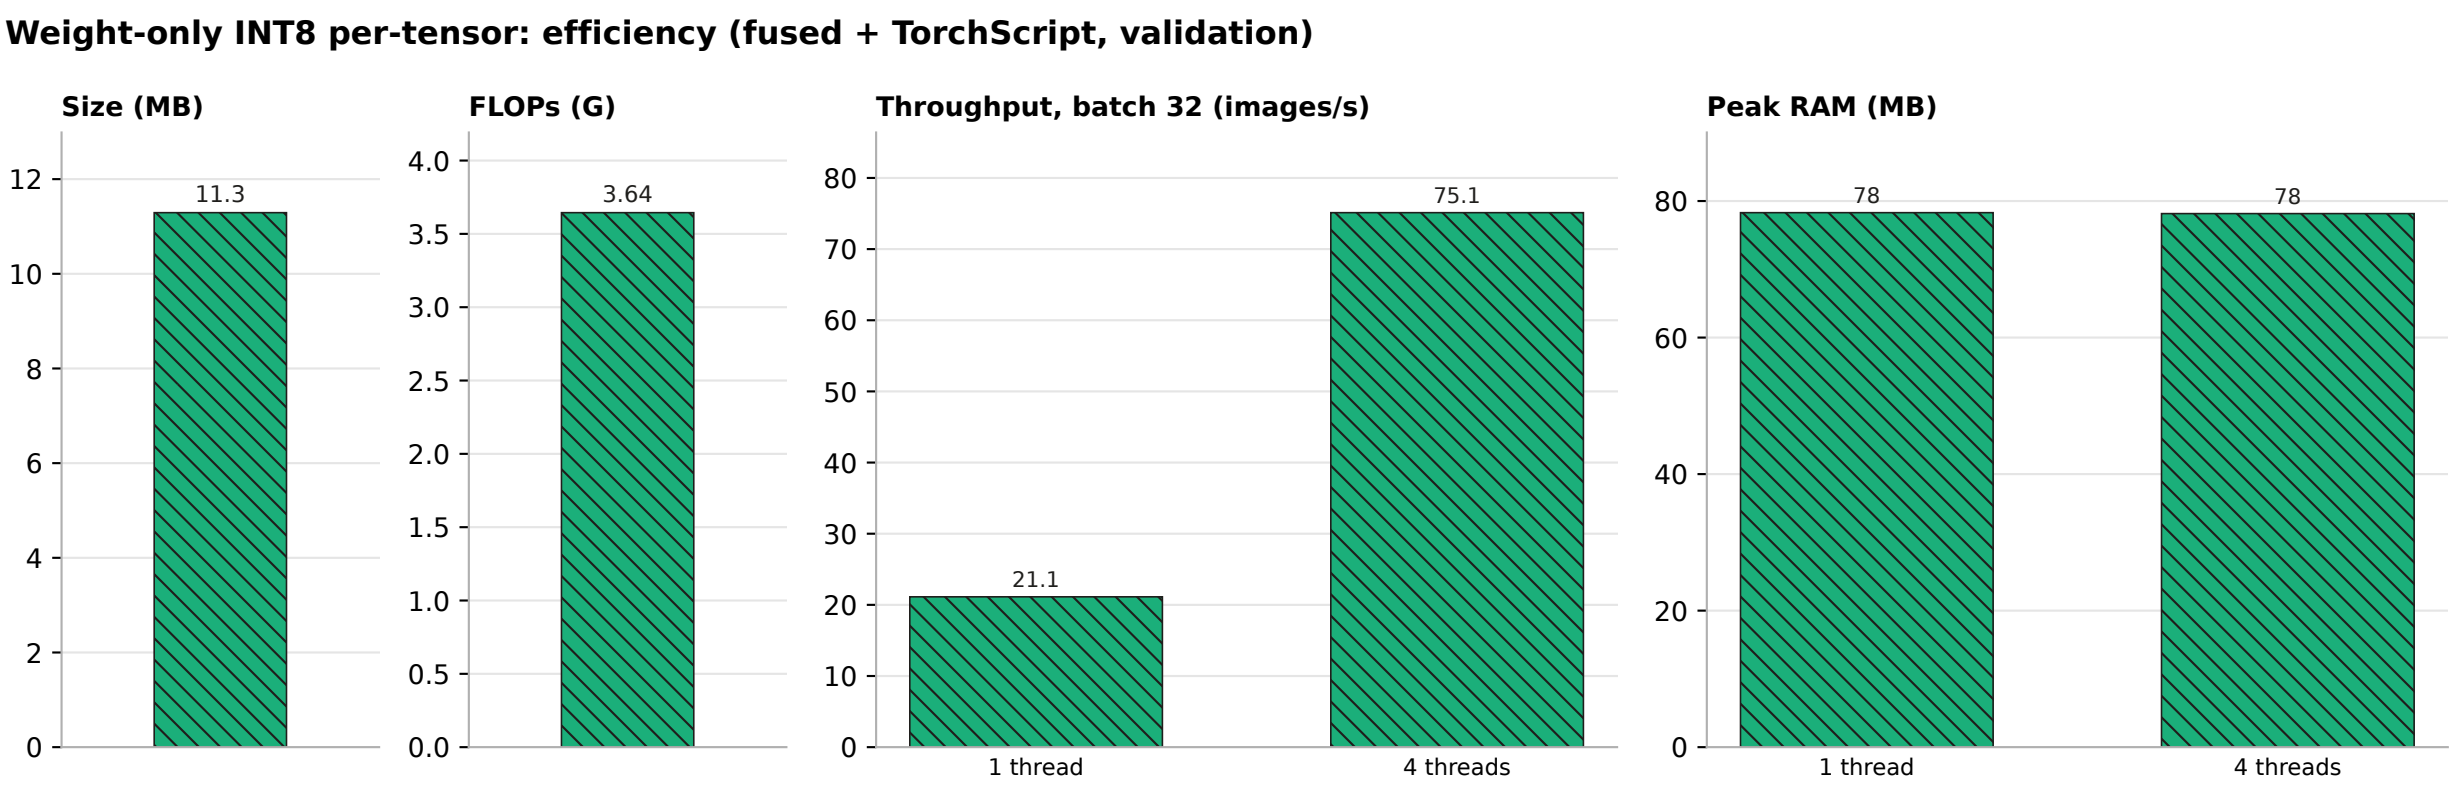
Weight-only INT8 per-tensor: efficiency (fused + TorchScript, validation)
Size (MB)
11.3
FLOPs (G)
3.64
Throughput, batch 32 (images/s)
21.1
75.1
1 thread
4 threads
Peak RAM (MB)
78
78
1 thread
4 threads