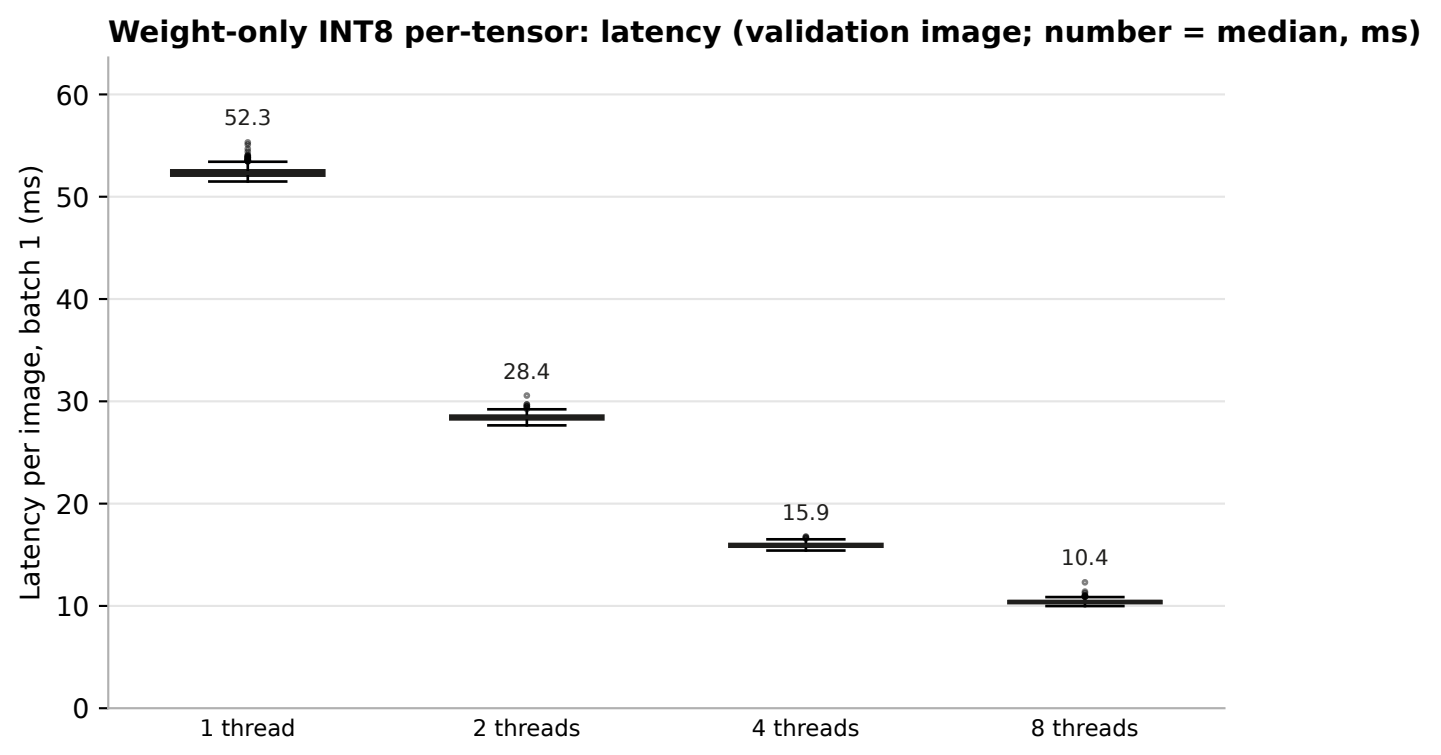
Weight-only INT8 per-tensor: latency (validation image; number = median, ms)
Latency per image, batch 1 (ms)
0
10
20
30
40
50
60
52.3
28.4
15.9
10.4
1 thread
2 threads
4 threads
8 threads
Each box: 900 timed runs (3 repeats x 300).