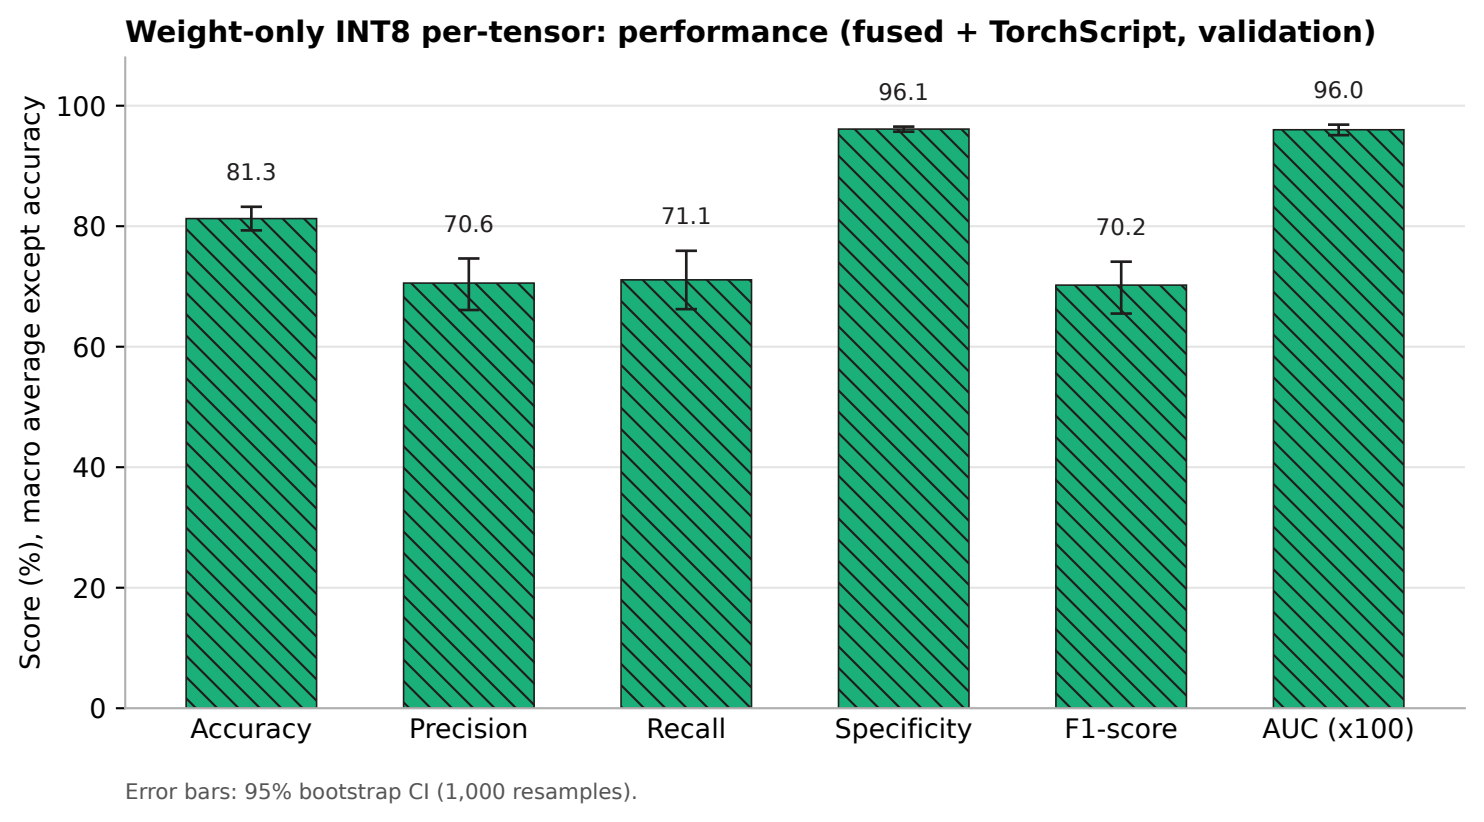
Weight-only INT8 per-tensor: performance (fused + TorchScript, validation)
Score (%), macro average except accuracy
0
20
40
60
80
100
81.3
70.6
71.1
96.1
70.2
96.0
Accuracy
Precision
Recall
Specificity
F1-score
AUC (x100)
Error bars: 95% bootstrap CI (1,000 resamples).