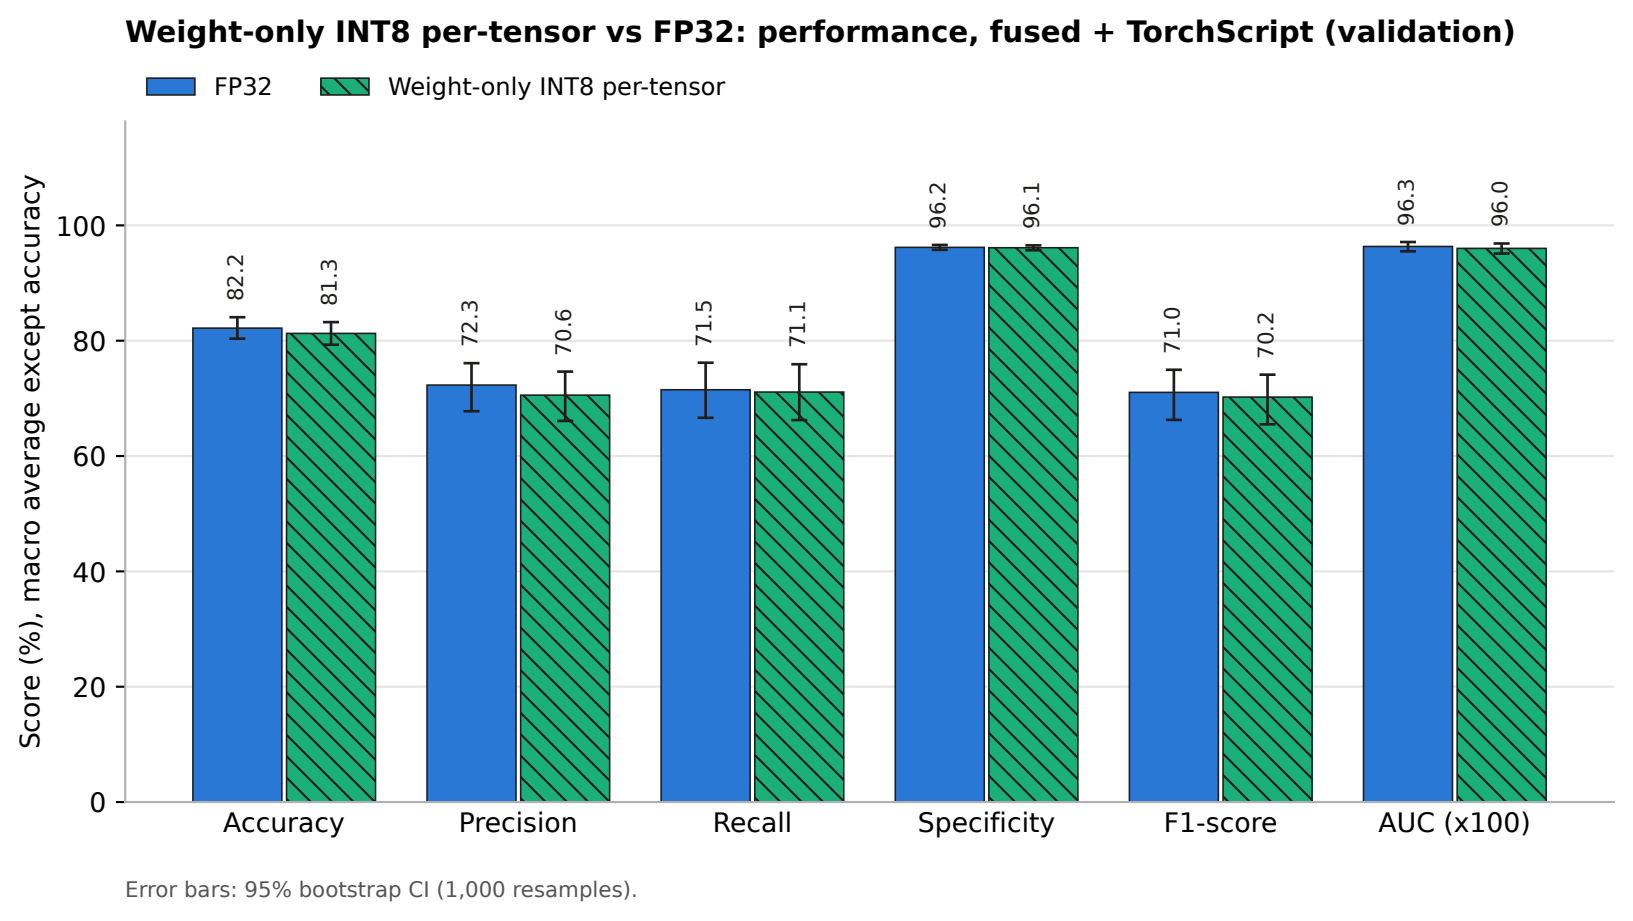
Weight-only INT8 per-tensor vs FP32: performance, fused + TorchScript (validation)
FP32
Weight-only INT8 per-tensor
Score (%), macro average except accuracy
0
20
40
60
80
100
82.2
81.3
72.3
70.6
71.5
71.1
96.2
96.1
71.0
70.2
96.3
96.0
Accuracy
Precision
Recall
Specificity
F1-score
AUC (x100)
Error bars: 95% bootstrap CI (1,000 resamples).