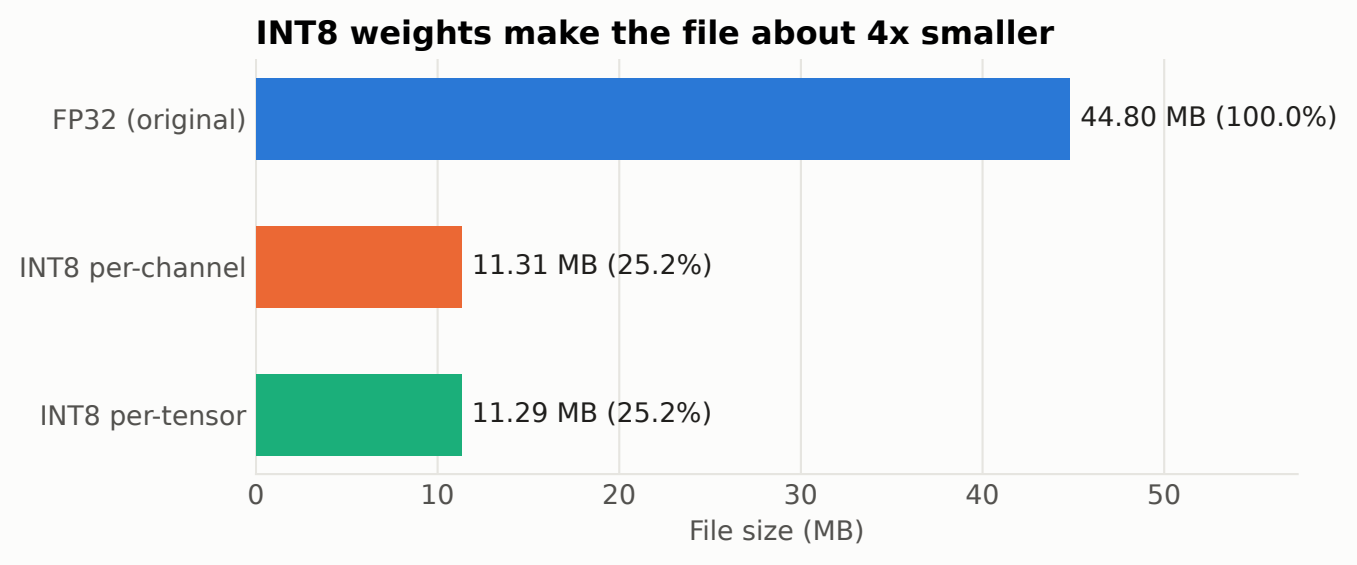
INT8 weights make the file about 4x smaller
FP32 (original)
INT8 per-channel
INT8 per-tensor
44.80 MB (100.0%)
11.31 MB (25.2%)
11.29 MB (25.2%)
0
10
20
30
40
50
File size (MB)